

Experimental Design and Data Analysis: Assignment 5

Andrew Bedard(2566978) & Simone van Gompel(2567525)
Group 19

June 4, 2015

Exercise 1

1

Using the data obtained from `nauseatable.txt`, we create a data frame consisting of two columns and 304 rows. One column contains an indicator, 0 for negative, 1 for positive, that a patient suffered from nausea, and the other column indicates the medicine the patient received.

2

By using the R command `xtabs(~medicine+nausea)` we are able to create a contingency table of our data so far. As we can see in figure 1

	No Nausea	Nausea
Chlorpromazine	100	52
Pentobarbital(100mg)	32	35
Pentobarbital(150mg)	48	37

Figure 1: Contingency table for all 3 drugs, and their effect

On this contingency table we are able to perform Pearson's Chi-squared test, performing this test we obtain a p-value of 0.0364, suggesting we reject the null hypothesis that there is no significant difference in the effect of the drugs.

3

We may also now perform a permutation test in order to determine whether the different medications work equally well against nausea. This is done by permuting the medicine labels, and once again using the Chi-squared test as our test statistic. We obtain the following histogram of values for the Chi-squared test, after permuting labels, and repeated 1000 times:




Figure 2: Histogram of Chi-squared values of permuted labels

By summing up all values obtained from this test which are greater than the value obtained by simply performing the Chi-squared test on the original data which is 6.62, then dividing by the number of samples performed, we create a p-value simulated in a boot-strap fashion. The resulting p-value is 0.033, which suggests that we reject the null hypothesis that all medicines work equally well.

4

- the p-value obtained from the Chi-square test for contingency tables was 0.0364
- the p-value obtained from the permutation test was 0.033

As a permutation test is a bootstrap type test, and the distribution of the test statistic is approximated by simulation we do see a variation in the p-values for each method, however for large replications for our permutation test on our data, this difference should not be significant enough to change our results. This is due to the fact that all our values of figure: 1 are greater than 1, and more than 80% of our values are greater than 5, which allows us to approximate the test statistic by the Chi-squared test.

Exercise 2

1

After reading in the data contained in `airpollution.txt`, we examine the relationship between all variables in figure 3

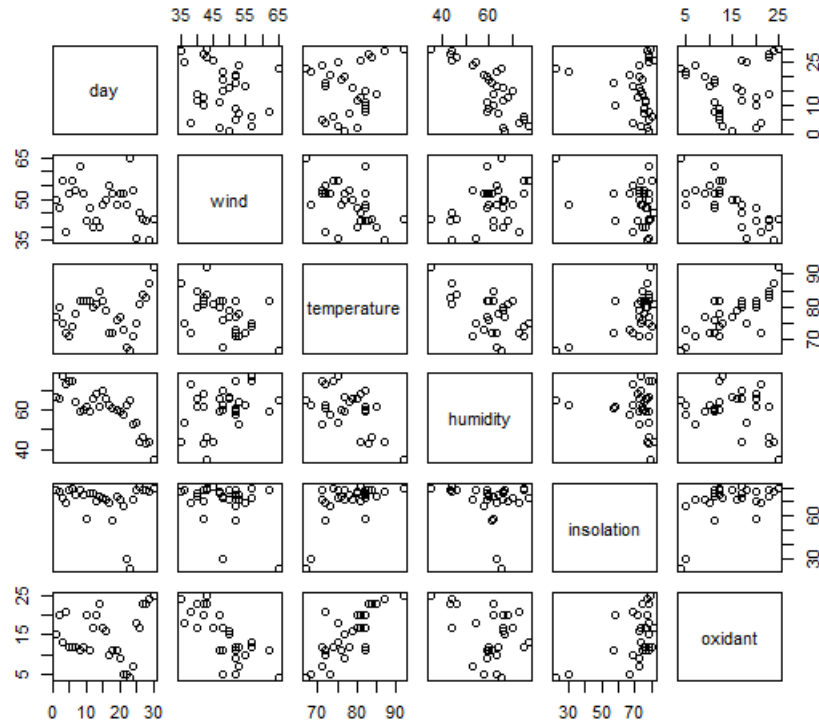


Figure 3: Pairplot of the airpollution data

As seen in above figure 3 there are many variables that have a visible linear relationship with oxidant, like temperature for example, this gives a strong indication that a linear model may be appropriate for this data.

2

With the response variable in this data being Oxidant, for each of the explanatory variables **day**, **wind**, **temperature**, **humidity**, **insolation**, we create a simple linear regression model, and using the method outline in R-Code section: 1.2 we stepwise extend the model by adding an additional explanatory variable based on increasing the determination coefficient of our model.

Table 1: Step-up method

Variable	R^2
Step1	
Oxidant~day	0.0109
Oxidant~wind	0.5863
Oxidant~temperature	0.5760
Oxidant~humidity	0.1240
Oxidant~insulation	0.2551
Step2	
Oxidant~wind+day	0.5988
Oxidant~wind+temperature	0.7773
Oxidant~wind+humidity	0.5913
Oxidant~wind+insulation	0.6613
Step3	
Oxidant~wind+temperature+day	0.7958
Oxidant~wind+temperature+humidity	0.7963
Oxidant~wind+temperature+insulation	0.7816
Step4	
Oxidant~wind+temperature+humidity+day	0.7974
Oxidant~wind+temperature+humidity+insulation	0.7979

As can be seen by Table: 1 after each step we add another variable to the model when that variable increases our determination coefficient R^2 , however the gains in the determination coefficient R^2 do not increase in any significant way after step 3, so this is where we stop. From this model we then obtain the estimates for our model parameters from a summary of our linear model:

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-16.60697	13.07154	-1.270	0.215
wind	-0.44620	0.08513	-5.241	1.78e-05 ***
temperature	0.60190	0.11764	5.117	2.47e-05 ***
humidity	0.09850	0.06316	1.559	0.131

Signif. codes:

0	***	0.001	**	0.01	*	0.05	.	0.1	1
---	-----	-------	----	------	---	------	---	-----	---

Figure 4: Summary of computed linear model

Our model is thus:

$$\text{Oxidant} = -16.6069 - 0.4462 * \text{wind} + 0.6019 * \text{temperature} + 0.0985 * \text{humidity} + \text{error}$$

Using a similar method to the previous section, we now stepwise decrease a model that includes all explanatory variables, and using a test of the form $H_0 : \beta_i = 0$ we select

variables to exclude from the model until we reject the null hypothesis for all explanatory variables still contained in the model.

Coefficients :

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-12.04010	21.20961	-0.568	0.57553	
day	-0.02997	0.13995	-0.214	0.83227	
wind	-0.44749	0.09103	-4.916	5.14e-05	***
temperature	0.55714	0.15347	3.630	0.00133	**
humidity	0.06818	0.13336	0.511	0.61384	
insolation	0.01822	0.05583	0.326	0.74694	

(a) Day has the highest p-value, we eliminate it from the model

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-15.49370	13.50647	-1.147	0.26219	
wind	-0.44291	0.08678	-5.104	2.85e-05	***
temperature	0.56933	0.13977	4.073	0.00041	***
humidity	0.09292	0.06535	1.422	0.16743	
insolation	0.02275	0.05067	0.449	0.65728	

(b) Insulation has the highest p-value, we eliminate it from the model

Figure 5: Steps 1 and 2 of step-down method

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-16.60697	13.07154	-1.270	0.215	
wind	-0.44620	0.08513	-5.241	1.78e-05	***
temperature	0.60190	0.11764	5.117	2.47e-05	***
humidity	0.09850	0.06316	1.559	0.131	

(a) Humidity has the highest p-value, we eliminate it from the model

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-5.20334	11.11810	-0.468	0.644	
wind	-0.42706	0.08645	-4.940	3.58e-05	***
temperature	0.52035	0.10813	4.812	5.05e-05	***

Signif. codes:

0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

(b) Summary of computed full linear model minus day, insulation and humidity, and all remaining variables have a p-value less than 0.05

Figure 6: Steps 3 and 4 of step-down method

Thus as illustrated in figure 5 and 6, at each step we remove a single explanatory variable that has a p-value greater than 0.05, or rather, a variable for which we fail to reject the null hypothesis that $\beta_i = 0$.

By this method we obtain the model:

$$Oxidant = -5.2033 - 0.4271 * wind + 0.5204 * temperature + error$$

4

The models that we obtained through the step up and step down methods are as follows

- For the step up model we obtained:

$$Oxidant = -16.6069 - 0.4462 * wind + 0.6019 * temperature + 0.0985 * humidity + error$$

- For the step down model we obtained:

$$Oxidant = -5.2033 - 0.4271 * wind + 0.5204 * temperature + error$$

The model that we would choose would be the one obtained through the step up method: $Oxidant = -16.6069 - 0.4462 * wind + 0.6019 * temperature + 0.0985 * humidity + error$. This is mainly due to the fact that our R^2 value indicates that the model does improve with the addition of humidity, as seen in part 2, even though we failed to reject the null hypothesis for this factor being 0 in the step-down method, thus our determination coefficient is 0.7963 as opposed to the step-down model which gives a determination coefficient of 0.7773.

5

We investigate the qqplot and the fitted plot of the residuals from the model we chose and obtained the following:




Figure 7: QQ and fitted plots of residuals

As seen in figure 7, both our Normal qq plot and fitted plot suggest that our residuals are normal, which would suggest that our linear model is in fact appropriate.

Exercise 3

1

The data from `genal2.txt` was read in and transformed into a data frame with two columns as outlined in the R-Code section 1.3

2

Making a box plot of the data frame created from the previous section we obtain the following:




Figure 8: Boxplot of 15 mutation probabilities

Clearly from the curved shape of the box plot figure:8, we can see that the dependence of y on mut is obviously not linear, but appears with a distinct curve.

3

Next we fit a multiple linear regression model for y on mut and mut^2 using the code outlined in R-Code Section 1.3. We obtain the following table:

Coefficients :									
	Estimate	Std. Error	t	value	Pr(> t)				
(Intercept)	0.60538	0.06317	9.583	3.94e-12	***				
mut	-23.78772	3.63372	-6.546	6.50e-08	***				
mut2	417.96154	44.16824	9.463	5.68e-12	***				
Signif. codes:									
0	***	0.001	**	0.01	*	0.05	.	0.1	1

Figure 9: Summary of multiple linear model for y on mut and mut^2

Which from the estimates we can conclude our model is should be:

$$Y = 0.60538 - 23.78772 * mut + 417.96154 * mut^2 + error$$

4

We now make a plot of our model as a function of the mutation probability seen in figure 10.




Figure 10: Plot of multiple linear regression model for values of *mut* between 0.004 and 0.08

By taking the value of *mut* at which the *y* value given is lowest, we obtain the optimal mutation probability, which we have calculated to be 0.0284

5

- parameters were estimated from our data, being $\beta_0 = 0.60538$, $\beta_1 = -23.78772$, $\beta_2 = 417.96154$
- Had we estimated the parameters from the data with the mutation probabilities defined as labels rather than numerical values we would have at estimated at least 16 parameters. This is because, as we saw in figure 8 the relationship between *y* and *mut* is clearly not linear, but looks quadratic in nature, thus we would need **at least** 1 of our 15 *mut* variables to be quadratic. To obtain the most accurate model, we would likely have estimated 30 parameters, that is, a mut_i and mut_i^2 for each label *i*.

Exercise 4

The dataset *expensescrime* is exolored. The response variable is *expend* and the rest are independent variables. In Fig11 the pairs plot of the data is shown.

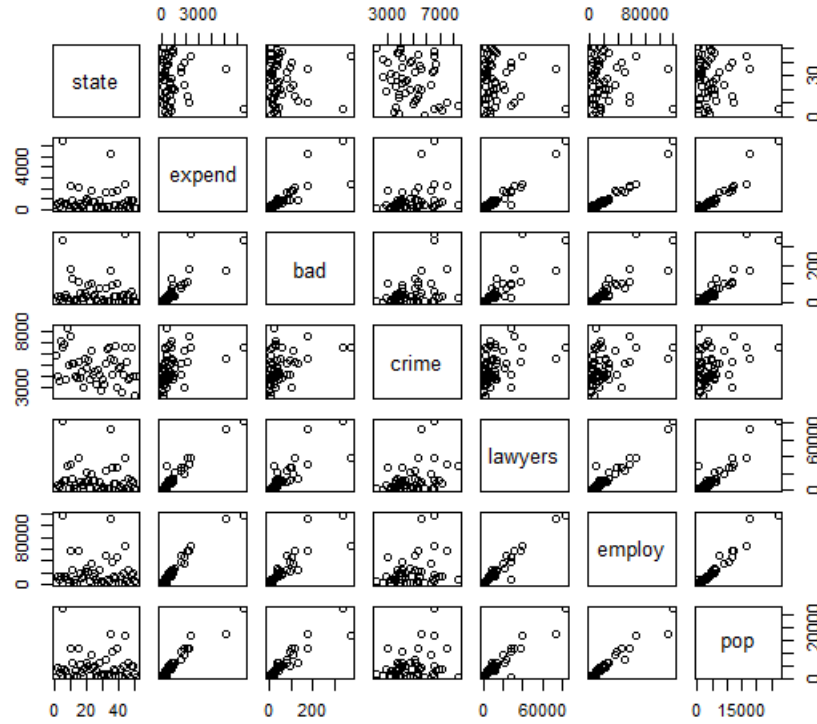


Figure 11: The pairs plot

From figure 11 we can see that many variables have a linear relationship with our response variable, suggesting that it would be appropriate to use a linear model. This figure also shows that many variables are co-linear. Indeed if we inspect pop and employ they have a very strictly linear relationship, and further from figure 12 we can see that their correlation is almost perfect at 0.97, which complicates our process and would suggest that we should not use both variables in our model.

	bad	crime	lawyers	employ	pop	expend
bad	1.00	0.37	0.83	0.87	0.92	0.83
crime	0.37	1.00	0.38	0.31	0.28	0.33
lawyers	0.83	0.38	1.00	0.97	0.93	0.97
employ	0.87	0.31	0.97	1.00	0.97	0.98
pop	0.92	0.28	0.93	0.97	1.00	0.95
expend	0.83	0.33	0.97	0.98	0.95	1.00

Figure 12: Correlation table

To create the linear regression model, we will utilize both the step-up and step-down methods to ensure that our model is as optimal as we can make it. The following is the result of the step-up method, similar as is illustrated in exercise 2

Table 2: Step-up method

Variable	R^2
Step1	
Expend~Pop	0.9073
Expend~Employ	0.954
Expend~Lawyers	0.9373
Expend~Crime	0.1119
Expend~Bad	0.6964
Step2	
Expend~Employ+Pop	0.9543
Expend~Employ+Lawyers	0.9632
Expend~Employ+Crime	0.9551
Expend~Employ+Bad	0.9551
Step3	
Expend~Employ+Lawyers+Pop	0.9637
Expend~Employ+Lawyers+Crime	0.9632
Expend~Employ+Lawyers+Bad	0.9639

Clearly from table 2 we can conclude that we do not achieve significant gains in our determination coefficient after adding employ and lawyers to our model, so we once again obtain our parameters from the following table:

Coefficients :					
	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-1.107e+02	4.257e+01	-2.600	0.01236	*
employ	2.971e-02	5.114e-03	5.810	4.89e-07	***
lawyers	2.686e-02	7.757e-03	3.463	0.00113	**
Signif. codes:					
0	***	0.001	**	0.01	*
				0.05	.
				0.1	
					1

Figure 13: Summary of computed linear model

Thus as a result of this step-up method, our model is:

$$Expend = -1.107e^2 + 2.971e^{-2} * employ + 2.686e^{-2} * lawyers + error$$

Again as illustrated in exercise 2, again we perform the step down method, eliminating variables that fail to reject the null hypothesis that their parameters are equal to 0.

Table 3: Step-down method

Variable	p
Step1	
Expend~Employ+Lawyers+Bad+Pop+Crime	$R^2 = 0.9675$
Employ	0.00354
Lawyers	0.00592
Bad	0.02719
Pop	0.03184
Crime	0.25534
Step2	
Expend~Employ+Lawyers+Bad+Pop	$R^2 = 0.9666$
Employ	0.00380
Lawyers	0.00106
Bad	0.05402
Pop	0.06012
Step3	
Expend~Employ+Lawyers+Bad	$R^2 = 0.9639$
Employ	1.2e-06
Lawyers	0.00147
Bad	0.34496
Step4	
Expend~Employ+Lawyers+Bad	$R^2 = 0.9632$
Employ	4.89e-07
Lawyers	0.00113

Both methods end up with the same formula, namely $\text{Expend} \sim \text{Employ} + \text{Lawyer}$. The results of the formula for Expend is as follows:

$$\text{Expend} = -1.107e^2 + 2.971e^{-2} * \text{Employ} + 2.686e^{-2} * \text{Lawyers} + \text{error}$$

Now that we have created a model, we should test if this result is reasonable, we have already concluded that some variables have collinearity, and we may still investigate whether there are influence points or outliers that may cause issues for us. Consider the following scatter plot of all variables left in our model:

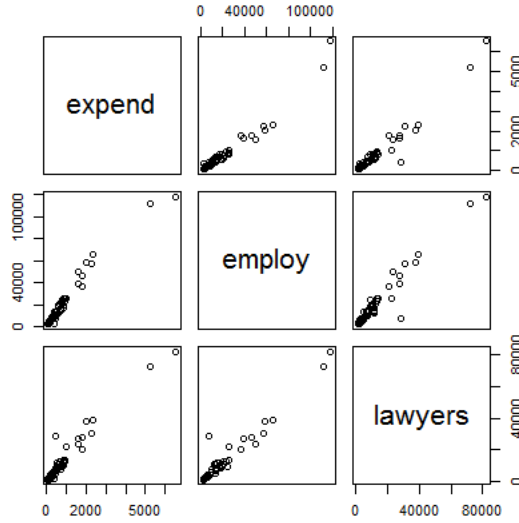


Figure 14: Scatter plots of all variables

It is very clear from figure 14 lawyers and employ may be collinear, and this is what we find in actuality, they are correlated by a factor of 0.9657, which is likely too high for them to both be in our model. So referring back to Table 2, we will conclude that our model shall only include employ, thus from the following table :

Coefficients :					
	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-1.167e+02	4.706e+01	-2.48	0.0166	*
employ	4.681e-02	1.469e-03	31.87	<2e-16	***

Signif. codes:									
0	***	0.001	**	0.01	*	0.05	.	0.1	1

Figure 15: Summary of computed linear model

we conclude that our model should now be:

$$Expend = -1.167e^2 - 4.681e^{-2} * employ + error$$

which yields only a slightly lower value for R^2 which is 0.954.

The influence points can be found by cook's distance, seen in Fig16. There is only significant data point that can influence the methods used.




Figure 16: The cook's distance

The residuals can be found in Fig17, which indicate that the data is not linearly correlated. None of the data in the plots is randomly scattered, This is backed up by Fig18

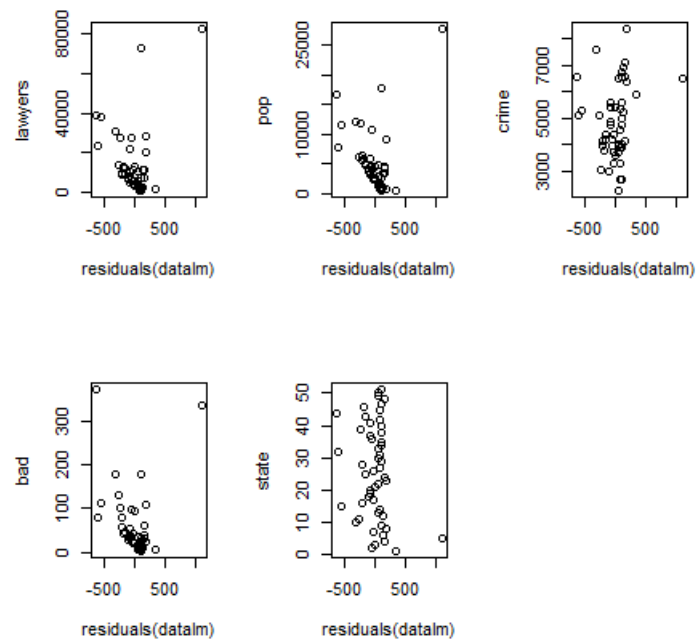


Figure 17: The residuals of the regression




Figure 18: QQ-plot of the residuals

Thus we conclude that our model is reasonable.

1 R-Code

1.1 Exercise 1

```

data = read.table("nauseatable.txt")

#1.1

rnames = rownames(data)
nausea = c()
medicin = c()
for (j in 1:length(data[,1])){
  for (i in 1:length(data[1,])){
    nausea = c(nausea, rep(i-1, data[j,i]))
    medicin = c(medicin, rep(rnames[j], data[j,i]))
  }
}
medicin <- factor(medicin)
nausea.frame=data.frame(nausea,medicin)

#1.2

#Create contingency table
cont_table = xtabs(~medicin+nausea)

#Contingency table has to be in matrix form for calculations

```

```

cont_matrix = matrix(c(100,52,32,35,48,37), byrow=TRUE, ncol=2 ,nrow=3,
                      dimnames=list(c(rownames(data)[1],
                                      rownames(data)[2],rownames(data)[3]),c("No_Nau",
                                      "Nausea")))

#chisquared test for contingency matrix
chisq.test(cont_matrix)

#1.3

teststat.obs = chisq.test(xtabs(~medicin+nausea))[[1]]

B = 1000
tstar=numeric(B)
for (i in 1:B)
{
  medstar = sample(medicin) ## permuting labels
  tstar[i] = chisq.test(xtabs(~nausea+medstar))[[1]]
}

hist(tstar)

#p-value

sum(tstar>teststat.obs)/B

#1.4

#chisquared test for contingency matrix with simulated p value
chisq.test(cont_matrix, simulate.p.value=TRUE)

```

1.2 Exercise 2

```

data = read.table("airpollution.txt")

#2.1
pairs(data)

#2.2

#create simple models, plot for reference
par(mfrow=c(2,3))

oxday = lm(oxidant~day, data=data)
plot(oxidant~day, data=data); abline(oxday)

```



```

oxwind = lm(oxidant~wind, data=data)
plot(oxidant~wind, data=data); abline(oxwind)

oxtemp = lm(oxidant~temperature, data=data)
plot(oxidant~temperature, data=data); abline(oxtemp)

oxhum = lm(oxidant~humidity, data=data)
plot(oxidant~humidity, data=data); abline(oxhum)

oxins = lm(oxidant~insolation, data=data)
plot(oxidant~insolation, data=data); abline(oxins)

summary(oxday)$r.squared
summary(oxwind)$r.squared
summary(oxtemp)$r.squared
summary(oxhum)$r.squared
summary(oxins)$r.squared

#wind has highest R squared, so add to model

summary(lm(oxidant~wind+day, data=data))$r.squared
summary(lm(oxidant~wind+temperature, data=data))$r.squared
summary(lm(oxidant~wind+humidity, data=data))$r.squared
summary(lm(oxidant~wind+insolation, data=data))$r.squared

#temperature has highest R squared, add to model

summary(lm(oxidant~wind+temperature+day, data=data))$r.squared
summary(lm(oxidant~wind+temperature+humidity, data=data))$r.squared
summary(lm(oxidant~wind+temperature+insolation, data=data))$r.squared

#adding humidity still increased R squared, add to model

summary(lm(oxidant~wind+temperature+humidity+day, data=data))$r.squared
summary(lm(oxidant~wind+temperature+humidity+insolation, data=data))$r.squared

#adding either day or insolation yields insignificant explanatory variables
#so we stop at previous step

summary(lm(oxidant~wind+temperature+humidity, data=data))

#summary gives coefficients of model in estimate column

#2.3

summary(lm(oxidant~day+wind+temperature+humidity+insolation, data=data))

```

```

#Try removing day, (highest p-value)
summary(lm(oxidant~wind+temperature+humidity+insolation , data=data))

#removing insolation , (highest p-value)
summary(lm(oxidant~wind+temperature+humidity , data=data))

#removing humidity, p-value too high still
summary(lm(oxidant~wind+temperature , data=data))

#2.4
#Present stepwise increase, less variables, but similar result

#2.5
step_up = lm(oxidant~wind+temperature+humidity , data=data)

par(mfrow=c(1,2))
qqnorm(residuals(step_up))
plot(fitted(step_up),residuals(step_up))

```

1.3 Exercise 3

```

data = read.table("genal2.txt")

#3.1
y = c()
mut = c()
# gives X0.005, so not useable
names = names(data)
length(data[1,])
length(data[,1])
for (i in 1:length(data[1,])){
  for (j in 1:length(data[,1])){
    y = c(y, data[j,i])
    mut = c(mut, 0.005*i)
  }
}
genal2frame=data.frame(y,mut)

#3.2
boxplot(data)

#3.3
genal2frame$mut2=genal2frame$mut^2
genal2lm=lm(y~mut+mut2, data=genal2frame)
summary(genal2lm)

```

#3.4

```
func1 = function(x) (0.605358-23.78772*x+417.96154*x^2)
```

```
B=1000
```

```
ypoints = numeric(B)
```

```
xpoints = seq(0.004,0.08,length=B)
```

```
for (i in 1:B){  
  ypoints[i] = func1(xpoints[i])  
}
```

```
plot(ypoints, xaxt = "n", xlab="mut_value", ylab="y_value", pch=20, cex=0.2, col="blue",  
axis(1, at=c(1,500,1000), labels=c(xpoints[1], round(xpoints[500], digits = 4), xpoints[1000]),  
#Optimal mutation value corresponds to xpoint where y is smallest  
round(xpoints[match(min(ypoints), ypoints)], digits=4))
```

1.4 Exercise 4

```
library('corrplot')
```

```
data = read.table("expensescrime.txt", header=T)
```

```
attach(data)
```

```
# looking at correlations,
```

```
# the correlation plot is on the last line of the program
```

```
pairs(data)
```

```
dataframe = data.frame(bad, crime, lawyers, employ, pop, expend)
```

```
round(cor(dataframe), 2)
```

```
# step-up algorithm
```

```
# examine r squared
```

```
summary(lm(expend~pop))$r.squared
```

```
summary(lm(expend~employ))$r.squared
```

```
summary(lm(expend~lawyers))$r.squared
```

```
summary(lm(expend~crime))$r.squared
```

```
summary(lm(expend~bad))$r.squared
```

```
summary(lm(expend~employ+pop))$r.squared
```

```
summary(lm(expend~employ+lawyers))$r.squared
```

```
summary(lm(expend~employ+crime))$r.squared
```

```
summary(lm(expend~employ+bad))$r.squared
```

```
summary(lm(expend~employ+lawyers+pop))$r.squared
```

```
summary(lm(expend~employ+lawyers+crime))$r.squared
```

```
summary(lm(expend~employ+lawyers+bad))$r.squared
```

```

summary(lm(expend~employ+lawyers))

# step-down algorithm
# same result as step-up so continue with this
summary(lm(expend~employ+lawyers+bad+pop+crime))
summary(lm(expend~employ+lawyers+bad+pop))
summary(lm(expend~employ+lawyers+bad))
summary(lm(expend~employ+lawyers))

# influence points are the ones that stand out in cooks distance
data1m = lm(expend~employ)
round(cooks.distance(data1m),2)
par(mfrow=c(1,1))
plot(cooks.distance(data1m))

# problems with colinearity, the corrplot is on the bottom
round(cor(data[,3:7]),2)

# looking at the different residuals
par(mfrow=c(1,1))
plot(residuals(data1m),employ)
par(mfrow=c(2,2))
plot(residuals(data1m),lawyers)
plot(residuals(data1m),pop)
plot(residuals(data1m),crime)
plot(residuals(data1m),bad)
par(mfrow=c(1,2))
plot(residuals(data1m),expend)
plot(residuals(data1m),fitted(data1m))
par(mfrow=c(1,1))
qqnorm(residuals(data1m))

pairs(~expend+employ+lawyers)
cor(employ,lawyers)

```